

Project Objectives

Rewritten 2026-07-31 following lecturer feedback: objectives follow the SMART approach (Specific, Measurable, Achievable, Relevant, Time-bound). Role-based access is a design feature, not an objective; it lives in the system design under the admin module.

Revised 2026-08-06: objectives 1 and 2 (repository, admin module) merged into a single objective, since the admin module's upload action already implements the repository's upload requirement, there was no separate build for each. Objective 5 (analytics dashboard) added; it reads data objective 3 already logs, so it adds no new data-collection requirement and does not extend the 9-week timeline.

Main Objective

To develop and empirically evaluate a centralized web platform that hosts university policy documents and answers natural-language policy questions through retrieval-augmented generation (RAG).

Specific Objectives

#	Objective (SMART)	Timeline	Measure of success
1	To develop, by the end of Week 2, a centralized web platform for university policy documents: PDF upload and storage into a single hosted repository.	Weeks 1–2	Platform hosts an initial corpus of ≥30 policy documents, and an administrator completes an add → replace → archive cycle entirely through the UI.
2	To integrate, by the end of Week 4, a RAG pipeline (vector embeddings + LLM, Chroma) that answers natural-language questions with citations to source policies.	Weeks 3–4	≥80% top-k retrieval precision on a benchmark set of 50 test questions.
3	To deploy the system, by the end of Week 6, to a test group of ≥20 users for a two-week period, logging every query (text, timestamp, retrieved documents, confidence score, unanswered queries) to the QueryLog table.	Weeks 5–6	Dataset of ≥200 logged queries from real users.
4	To analyze the collected data, by the end of Week 9, benchmarking retrieval accuracy and identifying the most frequently queried and most ambiguous policies.	Weeks 7–9	Evaluation report with accuracy metrics and policy-ambiguity charts for university leadership.
5	To surface, by the end of Week 6, the data objective 3 already logs through a live admin analytics dashboard: query volume over time, answered-vs-unanswered rate, and most-queried and most-ambiguous policies.	Week 6 (concurrent with objective 3)	Dashboard renders all four metrics live from the QueryLog table objective 3 already populates, no new logging beyond it.

Total duration: 9 weeks. Objective 5 runs concurrently with objective 3, it reads the same data objective 3 is already committed to collecting, so it does not extend the timeline.

Notes

- Target numbers (30 documents, 20 users, 200 queries, 80% precision) are working defaults; adjust to cohort size and supervisor guidance.
- Objective 5 is deliberately scoped to read-only reporting on existing data. If it turns out to need a metric objective 3 doesn't already log, that's new scope for objective 3, not a silent expansion of objective 5.